

LambdaGen

A GPU Code Generator Powered by Recursion Schemes

Dániel Berényi
Wigner Research Centre for Physics
GPU Lab

András Leitereg, Gábor Lehel
Eötvös University

Budapest Haskell Hackathon 2017



Wigner Research Centre for Physics:

- One of the largest research institute of the Hungarian Academy of Sciences
- Member of many important international collaborations: CERN LHC (particle physics), LIGO/VIRGO (gravitational waves), ESA (Rosetta mission), ITER, Jet (fusion experiments)...
- Wigner Datacenter
largest off-site compute infrastructure of the CERN

Introduction

Wigner GPU Lab:

- Research and support group at the Wigner Institute providing
 - Computational resources:
small GPU cluster and development machines from all vendors
 - Developer's assistance:
Help researchers with programming, dev tools, recommendations
Dissemination: annual [GPU Day](#), [Lectures on Modern Scientific Programming](#)
- Research/Develop scalable, generic simulations and visualizations
- Seek and Evaluate new, emerging technologies, participate in the development of existing ones
we're members of the Khronos OpenCL Advisory Panel



Introduction

We are always looking for tools that help to map abstract mathematical constructs to hardware

Functional programming provides a vast pool of such tools

Structures

Structures Trees

Structures

Trees

Hierarchies

Structures
Trees
Hierarchies
Recursion

Physicists have lots of nice and compact equations
e.g. Maxwell's eqs:

$$\partial_{[\alpha} F_{\beta\gamma]} = 0 \qquad \partial_{\alpha} F^{\alpha\beta} = \mu_0 J^{\beta}$$

that hide a tremendous amount of assumptions,
constructs and abbreviations...

Tree structures


Many times the problems are formulated at the symbolic level

But require low-level optimal codes on the hardware

$$\partial_{[\alpha} F_{\beta\gamma]} = 0$$

$$\partial_{\alpha} F^{\alpha\beta} = \mu_0 J^{\beta}$$

How to get there?



Tree structures

Usually the path involves many steps and tools:

Symbolic math expressions


Dense or sparse linear algebraic expressions

High-level programming terms

Low-level programming terms

$$\partial_{[\alpha} F_{\beta\gamma]} = 0 \quad \partial_{\alpha} F^{\alpha\beta} = \mu_0 J^{\beta}$$

How to get there?



Tree structures

Usually the path involves many steps and tools:

Symbolic math expressions

Dense or sparse linear algebraic expressions

High-level programming terms

Low-level programming terms

Symbolic transformations,
differentiation,
integration,
simplification,
substitution



Tree structures

Usually the path involves many steps and tools:


Symbolic math expressions

Dense or sparse linear algebraic expressions

High-level programming terms

Low-level programming terms

Numerical solver choices based on the structures of equations, matrix-tensor manipulations, decomposition into well-known primitive operations



Tree structures

Usually the path involves many steps and tools:

Symbolic math expressions

Dense or sparse linear algebraic expressions

High-level programming terms

Low-level programming terms

Types

Higher-order functions

Fusion-Fission

Specialization of generic algorithms



Tree structures

Usually the path involves many steps and tools:

Symbolic math expressions

Dense or sparse linear algebraic expressions

High-level programming terms

Low-level programming terms

memory layout & usage,
temporaries,
scheduling...



So there are trees everywhere:

- Expression trees (symbolic or program)
- Abstract Syntax Trees
(often code generation and code transformations are needed)
- Nested generic algorithms
- Nested data structures
- Hardware hierarchies

Tree structures

Hardware hierarchies:
Hierarchical parallelism:

Computing center
Clusters of computers
Multiple devices (CPU, GPU, FPGA)
Multiple execution units
Groups of threads



Tree structures

Hardware hierarchies:

Memory hierarchy:

Speed



Storage (disk, tape)

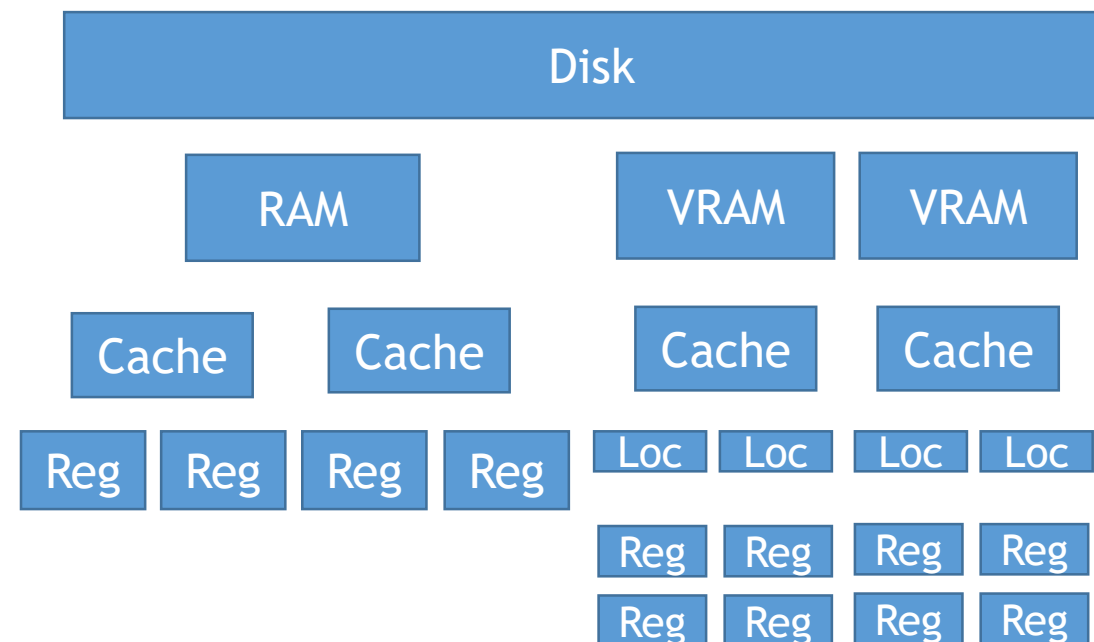
Device memory (RAM, VRAM)

Caches

Memory shared by threads

Registers

Size



So how to deal with all these trees?

All these trees are made from the same building blocks (branches and leaves) repeatedly under each other

→ trees are recursive structures

Primitive Recursion

```
data Tree = Empty | Branch Tree Tree
```

```
depth :: Tree -> Int
```

```
depth (Empty) = 0
```

```
depth (Branch l r) = 1 + max (depth l) (depth r)
```

```
d0 = Empty
```

```
d1 = Branch Empty Empty
```

```
d2 = Branch (Branch Empty Empty) Empty
```

```
main = print $ depth d2
```

Primitive Recursion

```
data Tree = Empty | Branch Tree Tree
```

```
depth :: Tree -> Int
```

```
depth (Empty) = 0
```

```
depth (Branch l r) = 1 + max (depth l) (depth r)
```

Primitive Recursion

The defined name also
appears on the rhs

```
d0 = Empty
```

```
d1 = Branch Empty Empty
```

```
d2 = Branch (Branch Empty Empty) Empty
```

```
main = print $ depth d2
```

Problems of this approach:

- Each recursive function had to be written with the particular datatype in mind (pattern match)
- Possible packing and unpacking of data transferred between levels is on the writer of the recursive function
- Same recursive logic might get implemented multiple times

Structured Recursion

```
{-# LANGUAGE DeriveFunctor #-}  
newtype Fix f = Fix (f (Fix f)) -- Fix point combinator  
  
unFix :: Fix f -> f (Fix f)  
unFix (Fix x) = x  
  
data TreeF r = Empty | Branch r r  
deriving (Show, Functor)  
  
type Tree = Fix TreeF  
  
cata :: Functor f => (f a -> a) -> Fix f -> a -- catamorphism  
cata alg = alg . fmap (cata alg) . unFix
```


Structured Recursion

```
{-# LANGUAGE DeriveFunctor #-}  
newtype Fix f = Fix (f (Fix f)) -- Fix point combinator
```

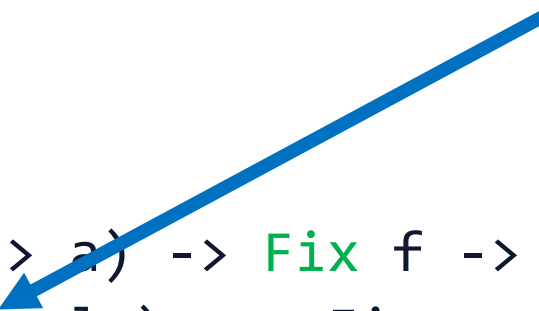
```
unFix :: Fix f -> f (Fix f)  
unFix (Fix x) = x
```

```
data TreeF r = Empty | Branch r r  
deriving (Show, Functor)
```

```
type Tree = Fix TreeF
```

```
cata :: Functor f => (f a -> a) -> Fix f -> a -- catamorphism  
cata alg = alg . fmap (cata alg) . unFix
```

We write the recursion logic once, this is the only place where the same name appears on the rhs too.



Structured Recursion

```
depth2 :: Tree -> Int
depth2 = cata alg
  where
    alg (Empty) = 0
    alg (Branch l r) = 1 + max l r

d0 = Fix $ Empty
d1 = Fix $ (Branch d0 d0)
d2 = Fix $ (Branch (Fix $ Branch d0 d0) d0)

main = print $ depth2 d2
```

Structured Recursion

```
depth2 :: Tree -> Int
```

```
depth2 = cata alg
```

```
  where
```

```
    alg (Empty) = 0
```

```
    alg (Branch l r) = 1 + max l r
```

Structured Recursion

This kind of traversal only depends on the evaluated results of the branches immediately below!

depth2 will not call itself directly!

```
d0 = Fix $ Empty
```

```
d1 = Fix $ (Branch d0 d0)
```

```
d2 = Fix $ (Branch (Fix $ Branch d0 d0) d0)
```

```
main = print $ depth2 d2
```

Structured Recursion

Why is this better?

- Separation of the recursion logic from the data type
- Parametrization over the recursive datatype and application logic
- More clearly expresses intent

Structured Recursion

Catamorphism:

Deconstruct a fixedpoint structure (**Functor** f)
into a summary value of type a:

```
cata :: Functor f => (f a -> a) -> Fix f -> a  
cata alg = alg . fmap (cata alg) . unFix
```

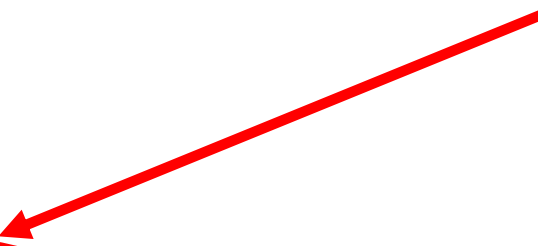
Structured Recursion

Catamorphism:

Deconstruct a fixedpoint structure (**Functor** *f*)
into a summary value of type *a*:

The deconstruction logic
(algebra) receives an
instance of the same functor
with the subresults

```
cata :: Functor f => (f a -> a) -> Fix f -> a  
cata alg = alg . fmap (cata alg) . unFix
```

A red arrow originates from the explanatory text on the right and points to the expression `(f a -> a)` in the Haskell code, which is circled in red. This expression represents the algebra function that takes a fixedpoint structure and returns a summary value.

Structured Recursion

Anamorphism:

Dual of catamorphism: constructs a structure from a seed

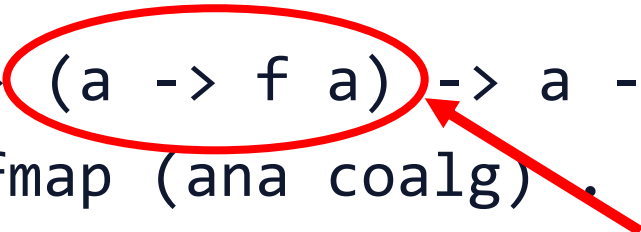
```
ana :: Functor f => (a -> f a) -> a -> Fix f
ana coalg = Fix . fmap (ana coalg) . coalg
```

Structured Recursion

Anamorphism:

Dual of catamorphism: constructs a structure from a seed

```
ana :: Functor f => (a -> f a) -> a -> Fix f
ana coalg = Fix . fmap (ana coalg) . coalg
```

A red circle highlights the expression `(a -> f a)` in the type signature of the `ana` function. A red arrow points from this circle to the `ana coalg` argument in the `fmap` function call within the `ana coalg` definition.

The construction logic (coalgebra) receives a seed value of type `a` and creates an instance of the same functor with the new seeds

Structured Recursion

Anamorphism:

Task: create a tree with given depth

```
plant :: Int -> Tree
plant = ana coalg
  where
    coalg 0 = Empty
    coalg n = Branch (n-1) (n-1)
```

Structured Recursion


Paramorphism (Generalization of the catamorphism):

Makes the original structure (before evaluation) available for the operation *beside* the evaluated sub-results:

```
para :: (Functor f) => (f (Fix f, a) -> a) -> Fix f -> a
para ralg = ralg . fmap fanout . unFix
  where --fanout :: Fix f -> (Fix f, b)
        fanout t = (t, para ralg t)
```

Paramorphism (Generalization of the catamorphism):

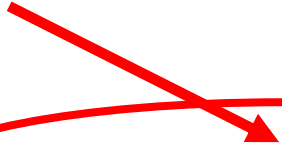
Same deconstruction like cata



```
para :: (Functor f) => (f (Fix f, a) -> a) -> Fix f -> a
para ralg = ralg . fmap fanout . unFix
  where --fanout :: Fix f -> (Fix f, b)
        fanout t = (t, para ralg t)
```

Paramorphism (Generalization of the catamorphism):

But the operation receives a structure of pairs:
where the left is the original subtree,
and the right is the subresult




```
para :: (Functor f) => (f (Fix f, a) -> a) -> Fix f -> a
para ralg = ralg . fmap fanout . unFix
  where --fanout :: Fix f -> (Fix f, b)
        fanout t = (t, para ralg t)
```

Structured Recursion

Paramorphism:

Task: count all half branches in the tree:



```
count :: Tree -> Int
```

```
count = para alg
```

```
  where
```

```
    alg Empty = 0
```

```
    alg (Branch (Fix Empty, _) (Fix Empty, _)) = 0
```

```
    alg (Branch (Fix (Branch _ _), n) (Fix Empty, _)) = 1+n
```


```
    alg (Branch (Fix Empty, _) (Fix (Branch _ _), n)) = n+1
```

```
    alg (Branch (Fix (Branch _ _), n) (Fix (Branch _ _), m)) = n+m
```

Structured Recursion - Use Cases

Use cases:

- Catamorphism
 - Decomposing structures level by level
 - Bottom-up traversal of a tree
 - Changing the expression structure type



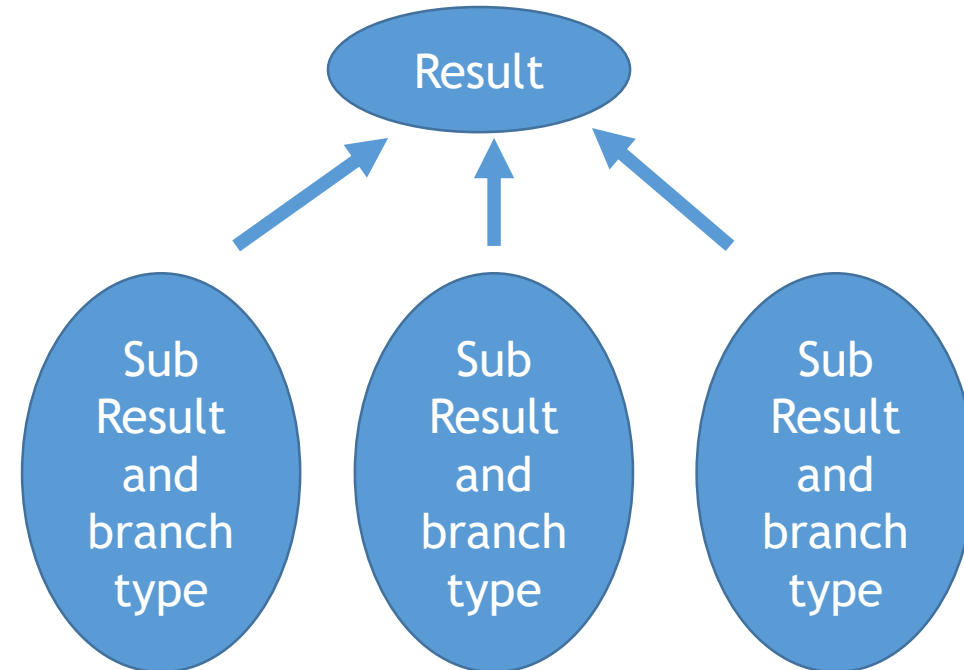
Evaluating, serializing expression trees
Calculating quantities bottom-up
Composition and combination identities



Structured Recursion - Use Cases

Use cases:

- Paramorphism
 - Decomposing structures level by level
 - Bottom-up traversal of a tree
 - Changing the expression structure type



Same things as cata, but ! we can depend on the structure of the subresults too!

Structured Recursion - Use Cases


Use cases:

- Anamorphism
 - Building structures level by level
 - Top-down traversal of a tree

De-serializing expression trees

Calculating quantities top-down

Create structure from seed



Structured Recursion

More examples for the case of expression tree manipulations: [see last year](#)

Structured Recursion

OK, so what is the hype for?

We just gave a fancy name for some frequent algorithm patterns...

Structured Recursion

OK, so what is the hype for?

We just gave a fancy name for some frequent algorithm patterns...

Not exactly...

Structured Recursion

Recursion schemes root in category theory, and form a nice hierarchy (khm☺) of more and more generic traversal patterns.

The theory gives also combination and fusion theorems that can be used to optimize composite expressions of recursion schemes.

Structured Recursion

A [collection](#) of recursion schemes is assembled from the literature by Edward Kmett and implemented in Haskell.

| Fold Schemes | Description | Unfold Schemes | Description |
|----------------|--|-----------------|--|
| Catamorphism | Consume structure level by level | Anamorphism | Create structure level by level |
| Paramorphism | Consume with primitive recursion | Apomorphism | Create structure, may stop and return with a branch or level |
| Zygomorphism | Consume with the aid of a helper function | | |
| Histomorphism | Consume, possibly multiple levels at once | Futumorphism | Create structure, possibly multiple levels at once |
| Prepromorphism | Consume, by repeatedly applying a natural transformation | Postpromorphism | Create, by repeatedly applying a natural transformation |

Structured Recursion

Combinations of generalized folds and unfolds are also available:

e.g.:

Hylomorphism = ana + cata

It can be used to build numerical quadrature algorithms

Mutumorphism: mutual recursive pair of functions

Can be used to create adaptive ordinary differential equation solvers

Switching from primitive recursion to structured recursion is just like switching from goto and explicit loops to packaged generic algorithms

They help seeing the pattern and solve complex problems.

Structured Recursion

An example outside of functional programming...

Neural Networks become new hype train in computing

Cheap, high-performance computing and large annotated datasets contributed to their success

One particularly important case of them are Recurrent Neural Networks, that are very good at natural language processing, handwriting, speech processing, recognition and generation.

Recurrent Neural Networks


- has a set of weights that describe a nonlinear transformation on an input
- and transform a sequential dataset by repeatedly applying itself on the inputs and the previous self state and/or output

Recurrent Neural Networks




Net


Recurrent Neural Networks




Recurrent Neural Networks




Recurrent Neural Networks



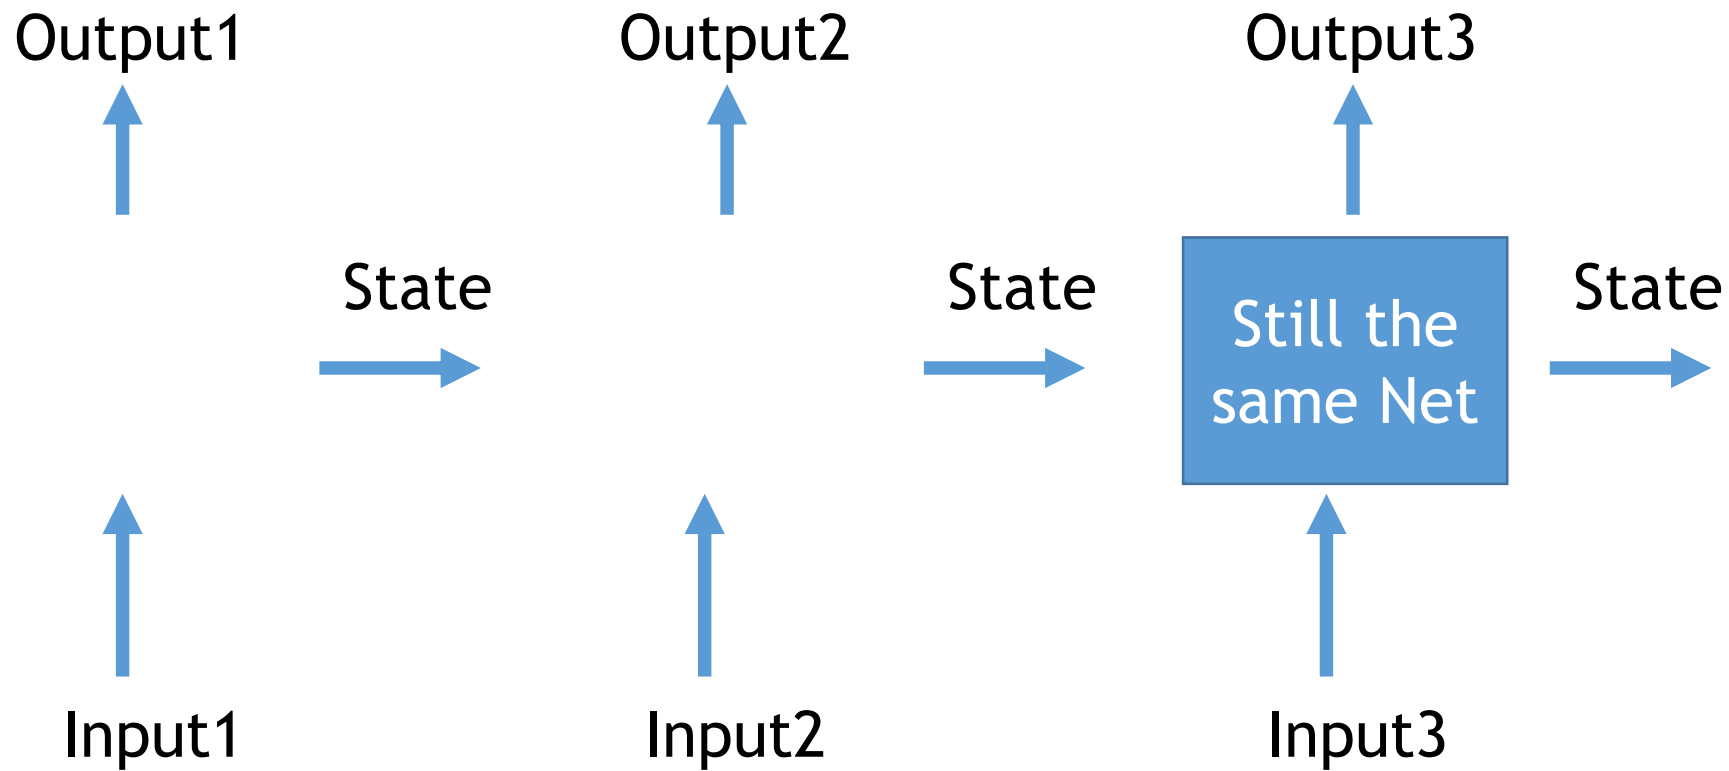
Recurrent Neural Networks



Recurrent Neural Networks



Recurrent Neural Networks




Recursive Neural Networks


Then it turned out that problems with more hierarchical structure need a different, more generic method:

Recursive neural networks, where the sequential evaluations are not a simple linear series, but form some sort of tree!


Recursive Neural Networks




Recursive Neural Networks




Recursive Neural Networks



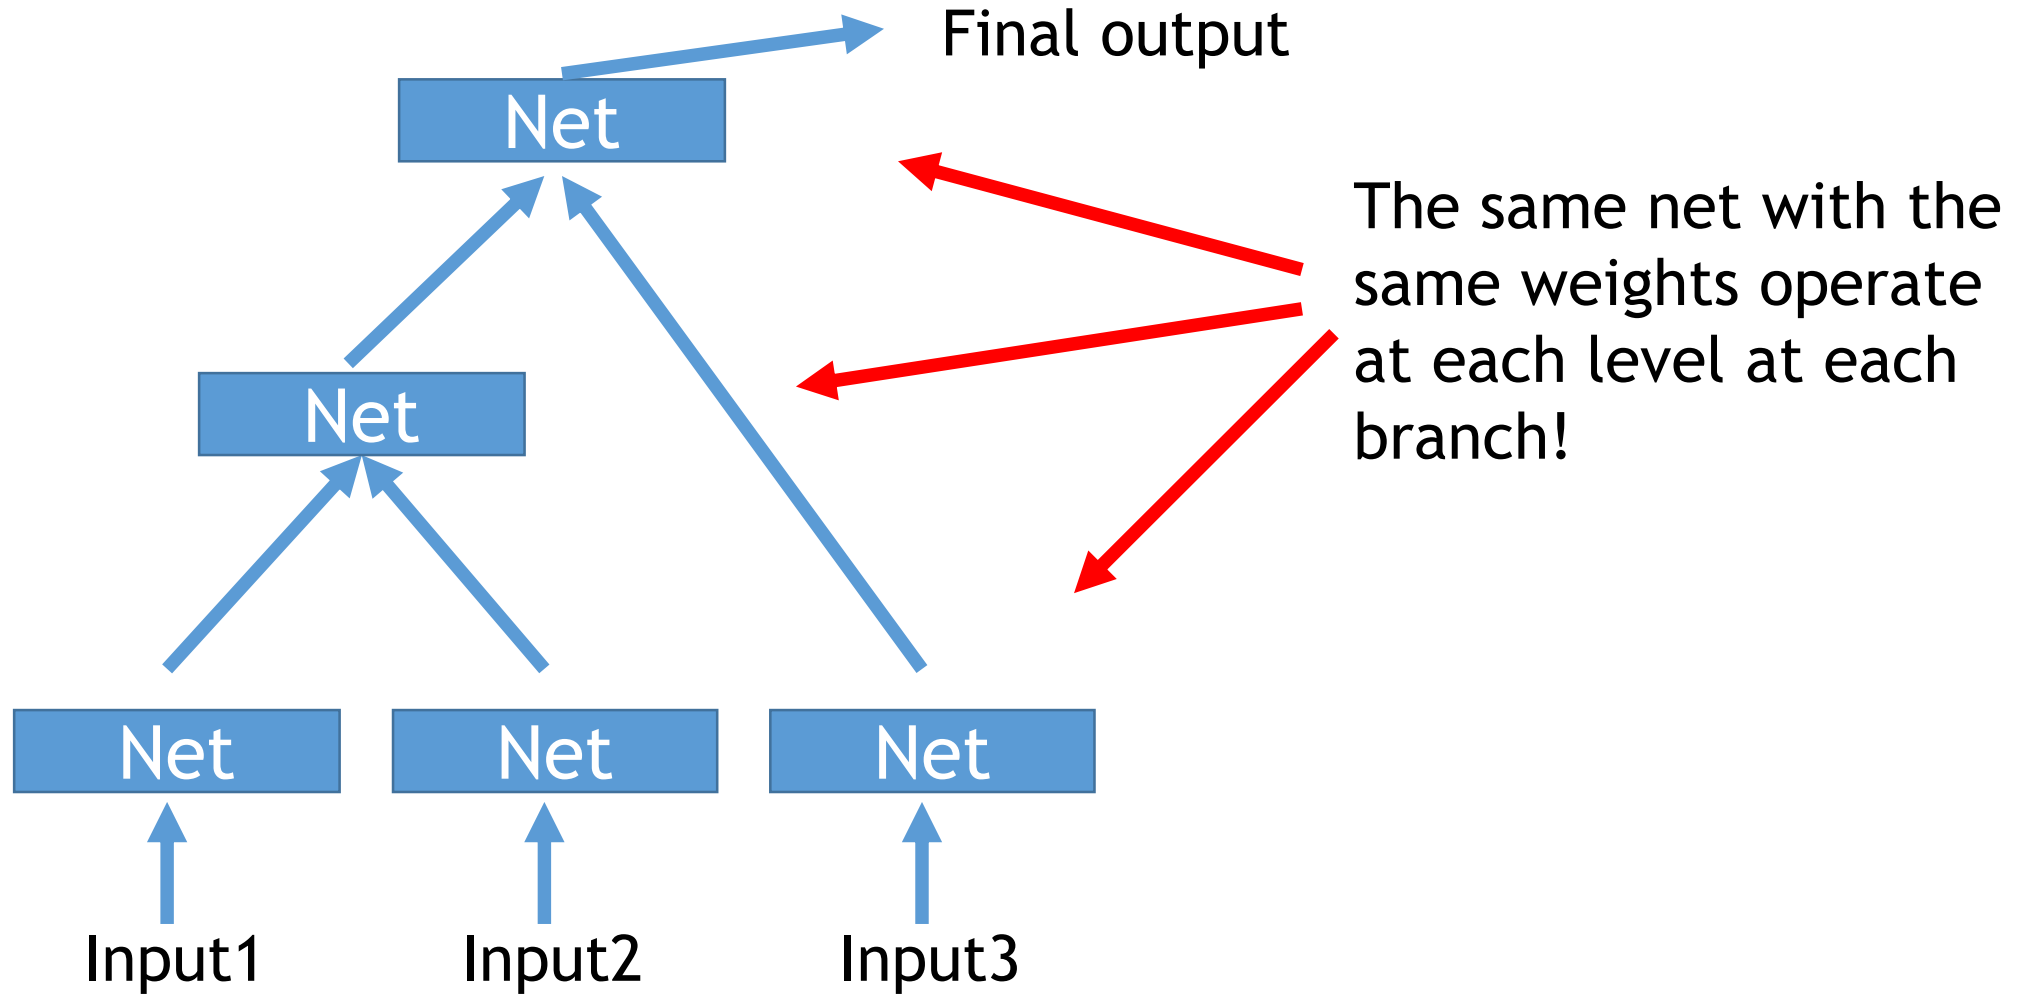
Recursive Neural Networks



Recursive Neural Networks



Recursive Neural Networks



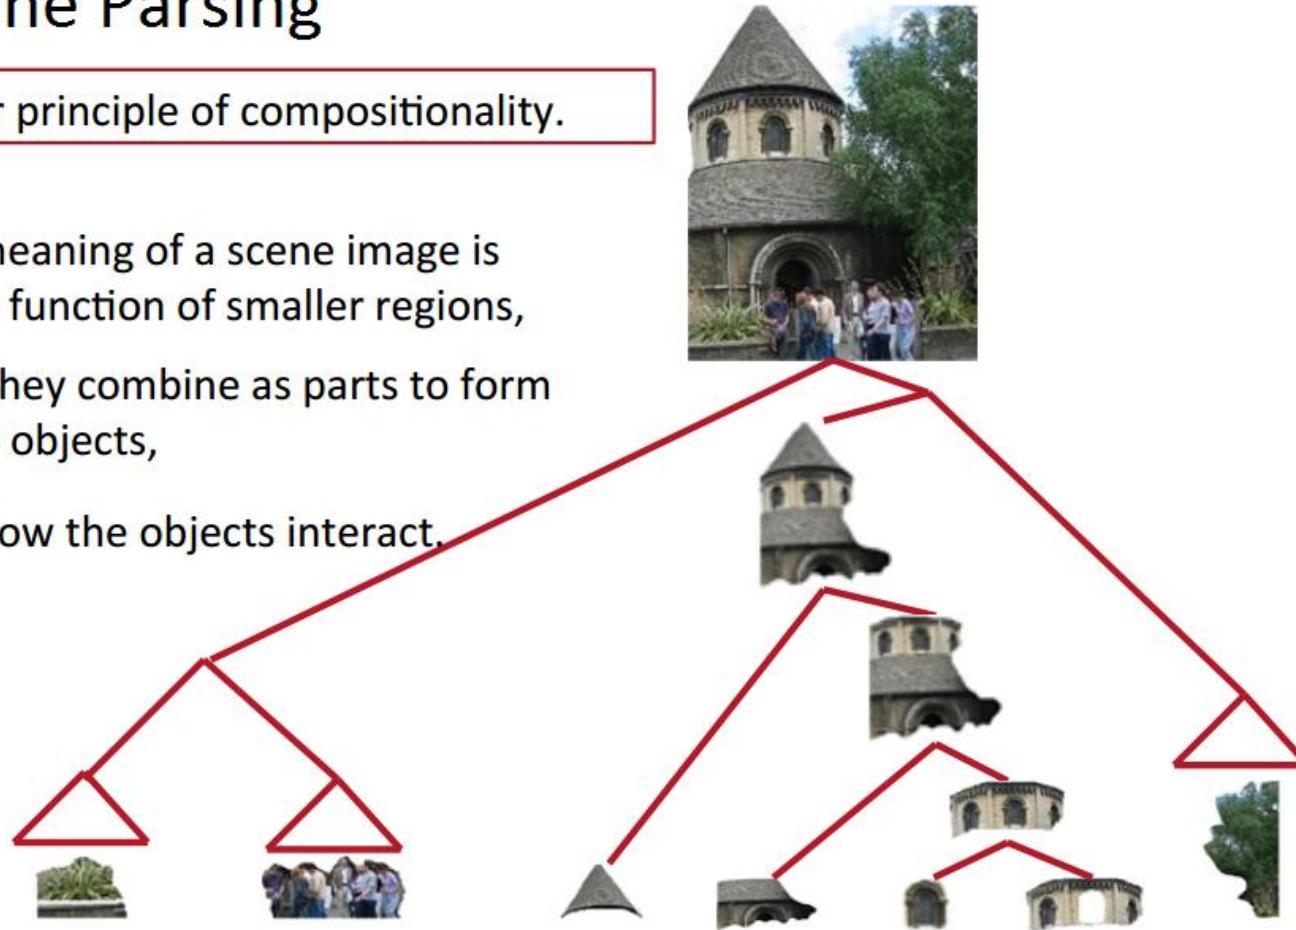
Recursive Neural Networks

Example:

Scene Parsing


Similar principle of compositionality.

The meaning of a scene image is also a function of smaller regions, how they combine as parts to form larger objects, and how the objects interact.



Socher, Richard & Chiung-Yu Lin, Cliff & Y. Ng, Andrew & Manning, Christopher. (2011).
Parsing Natural Scenes and Natural Language with Recursive Neural Networks.
Proceedings of the 28th International Conference on Machine Learning, ICML 2011. 129-136. 7/30/2017

Recursive Neural Networks



Same operation applied every time in a tree?

Does this sound familiar?

Christopher Olah pointed out [on his blog](#) that different neural network types closely correspond to functional programming primitives.

In particular:

- Recurrent Neural Networks are folds / unfolds
- Recursive Neural Nets are cata / anamorphisms

- Quite a lot of recursion schemes were described by the theoretical literature.
- Many of them are known to be useful and good to know about
- Recognizing them in different settings can make it easier to reason about and solve complex problems

Sample codes: github.com/u235axe/HaskellHacks

Erik Meijer, J. Hughes, M.M. Fokkinga, Ross Paterson

[Functional Programming with Bananas, Lenses, Envelopes and Barbed Wire](#)

Ralf Hinze, Nicolas Wu, Jeremy Gibbons

[Unifying Structured Recursion Schemes](#)

Tim Willams' [talk](#)

Edward Kmett's [blog](#) posts

Patrick Thomson's [blog](#) posts

Bartosz Milewski's [blog](#) posts

Recursive Neural Networks

Pollack, J. B. Recursive distributed representations. Artificial Intelligence Vol 46 (1990)

Bottou, L. arXiv.1102.1808. 401 (2011)

Frasconi, P., et. al. A general framework for adaptive processing of data structures. IEEE Transactions on Neural Networks , (1998).

Recursion Schemes - in real life

Next:

- András will show how to make a parallelizing compiler and GPU code generator by composing cata-, ana- and paramorphism

LambdaGen automates parts of the common scientific development pipeline:

- Symbolic math expressions
- Linear algebraic expressions
- High-level programming terms
- Low-level programming terms

Goal: minimize dev time + run time

$$\partial_{[\alpha} F_{\beta\gamma]} = 0 \quad \partial_{\alpha} F^{\alpha\beta} = \mu_0 J^{\beta}$$

How to get there?



LambaGen consists of

- A Haskell EDSL to define the computation
 - Functional primitives (map, reduce, zip, lambda calculus)
- A set of analysis and transformation steps
 - The primitives form an expression *tree*
 - all the previous theory applies
- GPU code generation ([SYCL](#) / [ComputeCpp](#))
- Automatic memory management (C++)

LambdaGen – Primitives

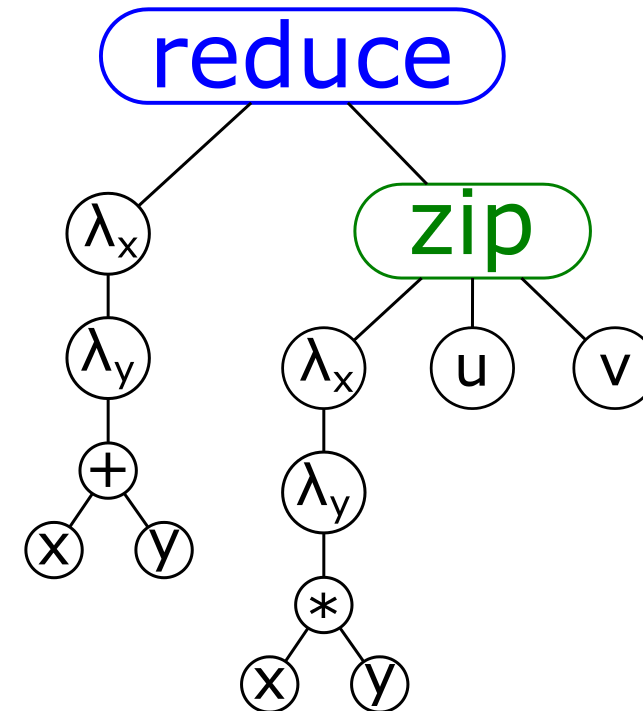

```
data ExprF a =
  Scalar      { getValue :: Double }
| Addition    { left :: a, right :: a }
| Multiplication { left :: a, right :: a }
| VectorView  { id :: String, dms :: [Int], strd :: [Int] }
| Apply       { lambda :: a, value :: a }
| Lambda      { varID :: String, varType :: Type, body :: a }
| Variable    { id :: String, tp :: Type }
| Map         { lambda :: a, vector :: a }
| Reduce      { lambda :: a, vector :: a }
| ZipWith     { lambda :: a, vector1 :: a, vector2 :: a }
deriving (Functor, Show)
```

LambdaGen – Example

$$\sum_i u_i v_i$$

reduce

```
(lam x (lam y  
  (add x y)))  
(zip  
  (lam x (lam y  
    (mul x y)))  
  u  
  v)
```



We can put annotations on the expression tree nodes by using

```
data Cofree f a = a :< f (Cofree f a)
    -- = a :< f (a :< f (Cofree f a))
```

Instead of

```
newtype Fix f = Fix (f (Fix f))
    -- = Fix (f (Fix (f (Fix f))))
```

Data.Functor.Foldable

- Fix, Cofree and the others are generalized into:

```
type family Base t :: * -> *
```

- Which for Fix means:

```
type instance Base (Fix f) = f
```

- And so the class of recursive types become:

```
class Functor (Base t) => Recursive t where
```


```
  project :: t -> Base t t -- unFix
```

```
  cata :: (Base t a -> a) -> t -> a
```


```
  cata alg = alg . fmap (cata alg) . project
```

... and countless other schemes ...

LambdaGen – Annotations



- The set of annotations changes while processing the expression tree

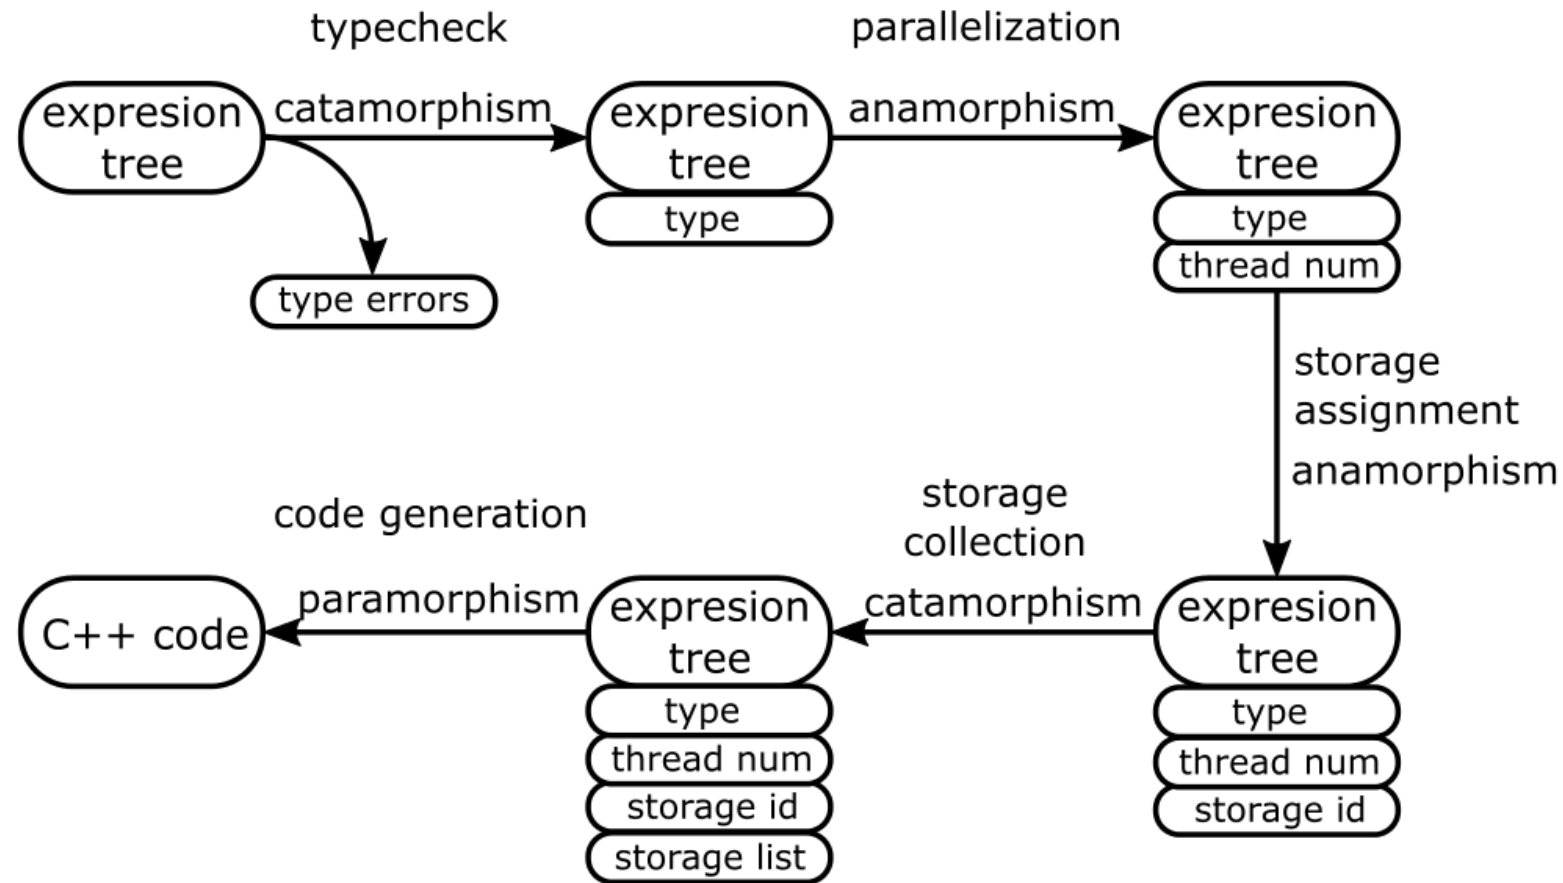


- A transformation must not depend on the exact set, but may require some annotations to be present


All this and much more is supported by the Vinyl extensible record package:

- The field set is a type level list \rightarrow extension is ' ' :
- Presence of fields can be constrained
 - `info1 \in fields \Rightarrow ...`
- Or even subset relationships
 - `subs \subseteq fields \Rightarrow ...`

LambdaGen – Overview



LambdaGen – Typecheck




```
type Algebra t a = Base t a -> a
```

```
type Error = String
```

```
type TypecheckT = Either Type [Error]
```


```
typecheckAlg :: Algebra (Cofree ExprF (R fields)) TypecheckT
```


LambdaGen – Annotations



The previous algebra would reduce the whole tree into a single type

But we'd like to make a chain of transformations




We should annotate the nodes with the partial results

LambdaGen – Algebra adapter

The `annotate` function takes a „reducing” algebra and creates a transforming or „annotating” algebra, that keeps the tree and adds some more info:

```
annotate :: Functor f =>
    Algebra (Cofree f (R old)) new ->
    Algebra (Cofree f (R old))
           (Cofree f (R (new ': old)))
```

LambdaGen – Parallelization



```
type CoAlgebra t a = a -> Base t a
data ParState = None | Started Int | Full Int
type ParData = (Int, Maybe Int)
```

```
parallelizerAlg :: TypecheckT ∈ fields => Int ->
    CoAlgebra (Cofree ExprF ParData)
              (Cofree ExprF (R fields), ParState)
```

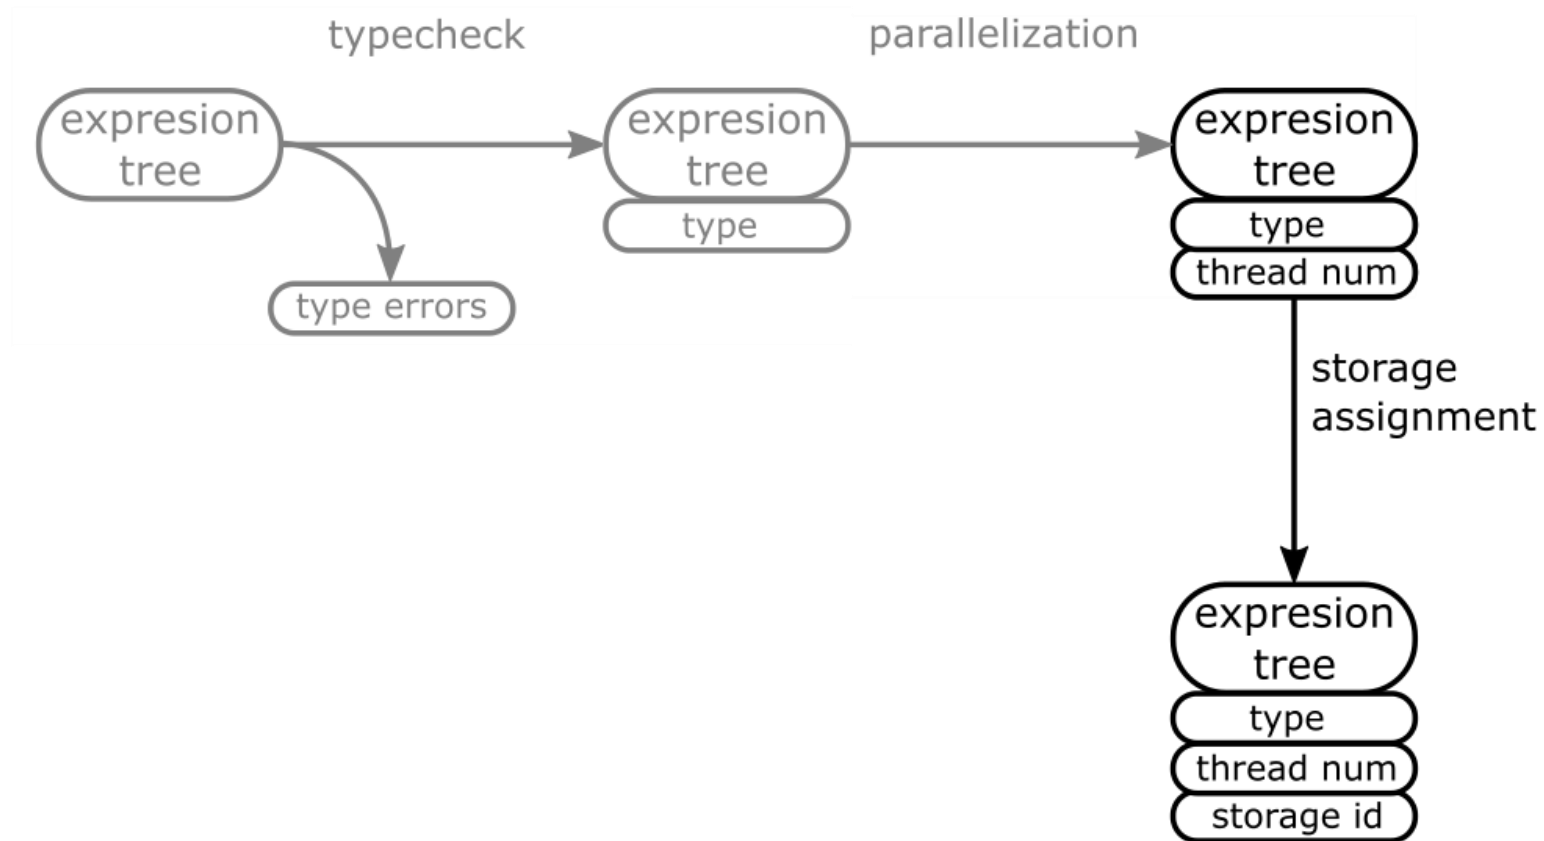
LambdaGen – CoAlgebra adapter

Just like before

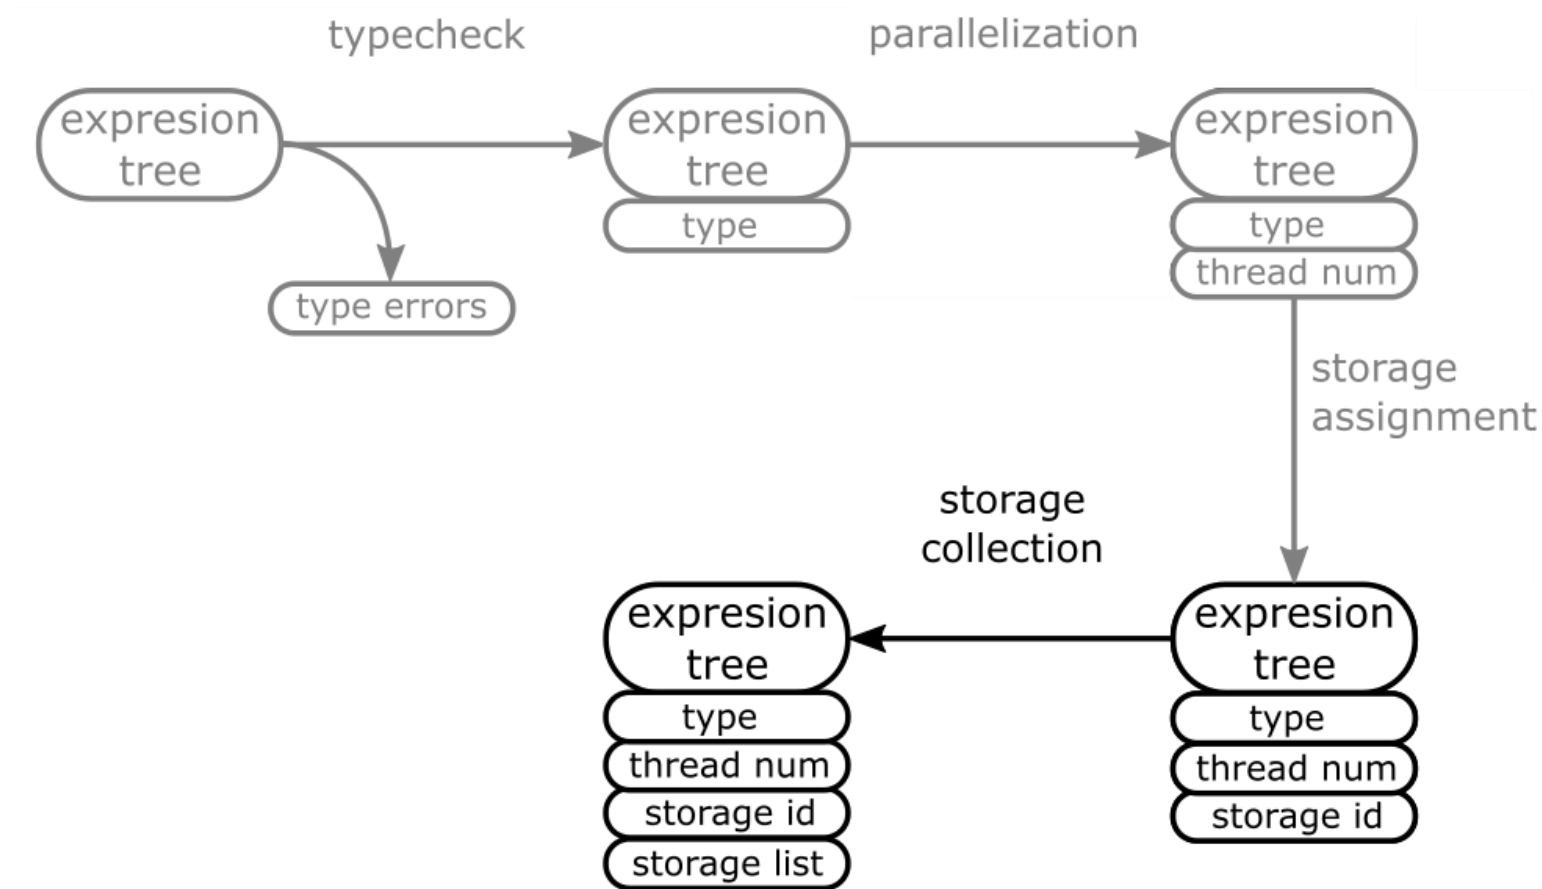
- we produce the new annotation
- the adapter appends it to the existing ones

```
annotateAna :: Functor f =>
    CoAlgebra (Cofree f new)
              (Cofree f (R old), seed) ->
    CoAlgebra (Cofree f (R (new ': old)))
              (Cofree f (R old), seed)
```

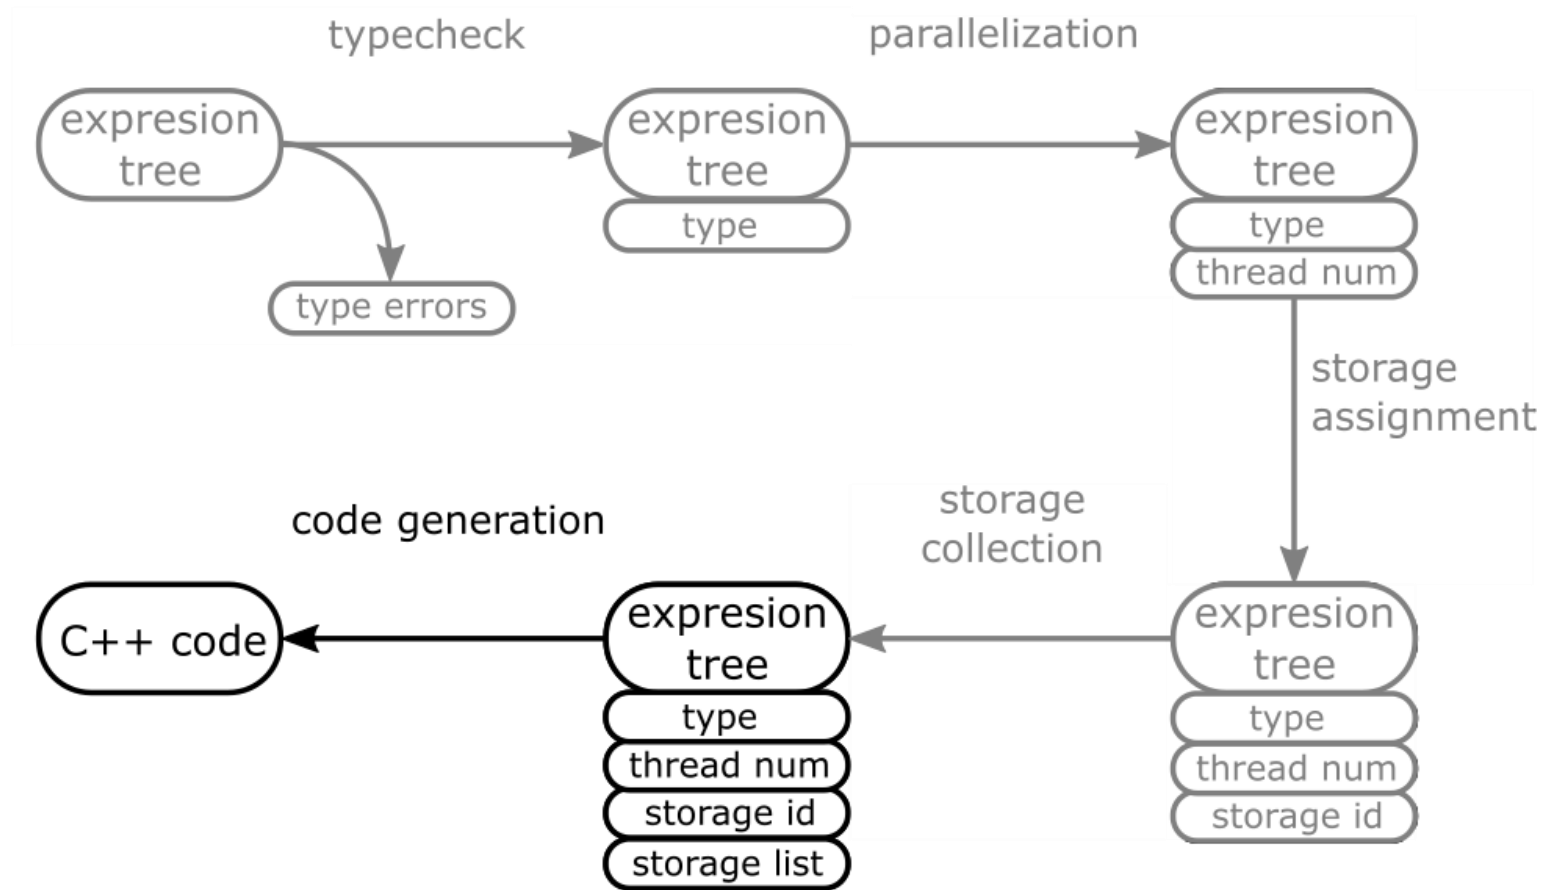
LambdaGen – Storage



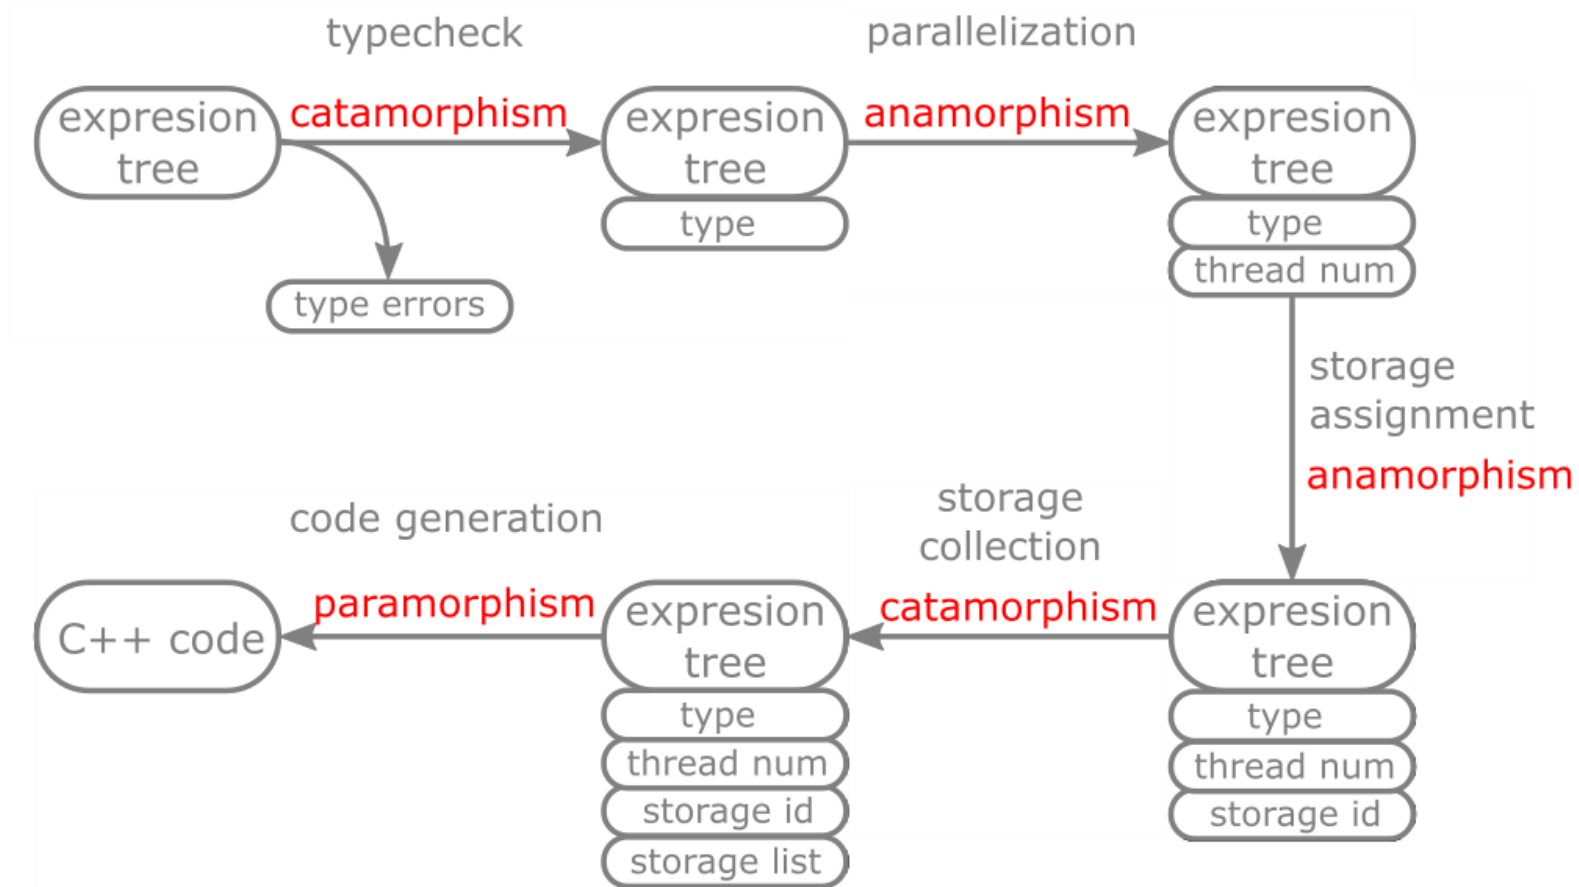
LambdaGen – Storage



LambdaGen – Code generation



LambdaGen – Recursion schemes



LambdaGen — Status

Tested on simple, but multidimensional linear algebra problems
(vector, matrix, tensor) (sums, products)

Works on CPU with `std::thread` based parallelism

Works on GPU, thanks to ComputeCpp's C++11 support

Code available on github:

github.com/leanil/LambdaGen

Further work:

- High-level optimization framework
 - Based on fusion/fission rules from theory
 - We'll need to pattern match and replace the expression tree
 - Can this be done with recursion schemes?
- Hierarchical parallelism
 - Local memory usage and vectorisation with HOF splitting and exchanging rules

Summary

Research at the **Wigner GPU Lab** aims to reduce the burden on scientists via bridging high level mathematical models down to hardware optimized codes with modern programming approaches

LambdaGen with its solely recursion schemes based extensible compiler logic is an important step in this direction

Follow or join the developments:

gpu.wigner.mta.hu

github.com/Wigner-GPU-Lab

github.com/leanil/LambdaGen